

Project #1 Tic-Tac-Toe  
CP468-Artificial Intelligence  
Ahmed Ibrahim  
3/31/2025

Ian Allan 210683230  
Justin Shirer 210269710  
Mobeen Akhtar 169048302  
Ahmed Sohail Butt  
Harveer Dhami 169032886  
Arshia Gole Sorkh 210747100  
Cagri Isilak 210764050  
Owen Schoeck 169020303  
Ryan Soomal 210370340

## Overview

- Implement the Minimax Algorithm: Develop a recursive decision-making process to determine optimal moves in an adversarial game environment.
- Enhance with Alpha-Beta Pruning: Improve efficiency by pruning suboptimal branches in the Minimax decision tree to reduce computational overhead.
- Integrate Gemini API: Use the Gemini API to play against advanced AI agents and evaluate performance under real-time decision-making.
- Compare Algorithm Performance: Measure and compare the execution time, node evaluations, and win rates of Minimax, Alpha-Beta Pruning, and Gemini agents.
- Analyze Scalability: Evaluate how the algorithms perform as game complexity increases, such as moving from 3x3 to larger grids (e.g., 5x5).

## Design

### Game Representation:

- The Tic-Tac-Toe board is modelled as a 3x3 grid using a 2D list.
- Each cell holds one of three possible values:
  - 'X' for Player 1
  - 'O' for Player 2
  - None or '-' for empty cell

### Turn-Based Agent Setup:

- The game alternates turns between two agents:
  - One uses Minimax or Alpha-Beta Pruning
  - The other can be controlled via the Gemini API
- Logic is structured to simulate real-time gameplay decisions.

### Heuristics:

- A basic utility evaluation is used:
  - +10 for a win
  - -10 for a loss
  - 0 for a draw
- Heuristics are only applied at terminal states since Tic-Tac-Toe is a simple game.

### Recursive Minimax Structure:

- Minimax uses a recursive depth-first approach to simulate every possible future game state.
- The algorithm chooses the move that maximizes or minimizes the player's score depending on whose turn it is.

### Alpha-Beta Pruning Optimization:

- Alpha-beta pruning reduces computational time by cutting off branches of the game tree that won't influence the final decision.
- This allows a deeper look ahead within the same time constraints.

### Modular Code Design:

- Code is organized into separate modules for:
  - Game logic
  - AI algorithms (Minimax & Alpha-Beta)
  - Gemini API communication
- This modular design improves maintainability and debugging.

## Algorithm Analysis

### Minimax Algorithm

The Minimax algorithm is a decision-making algorithm used in turn-based, perfect-information games (like Tic-Tac-Toe, Chess, etc.). It explores all possible future moves and chooses the optimal move for the maximizing player while assuming that the opponent plays optimally.

#### **Pseudocode**

```
function MINIMAX(board, player, maximizingPlayer):
    if game_over(board):
        return get_utility(board, maximizingPlayer)

    if player == maximizingPlayer:
        best_score = -∞ # Maximizing player wants the highest score
        for move in get_possible_moves(board):
            new_board = apply_move(board, move, player)
            score = MINIMAX(new_board, switch_player(player), maximizingPlayer)
            best_score = max(best_score, score)
        return best_score
    else:
        best_score = +∞ # Minimizing player wants the lowest score
        for move in get_possible_moves(board):
            new_board = apply_move(board, move, player)
            score = MINIMAX(new_board, switch_player(player), maximizingPlayer)
            best_score = min(best_score, score)
        return best_score
```

#### **Logic**

1. Base Case: If the game is over, return the utility value:

Win → +1 (for maximizing player)

Loss → -1 (for maximizing player)

Draw → 0

2. Recursive Case:

- If it's the maximizing player's turn, find the move that gives the highest possible score.
- If it's the minimizing player's turn, find the move that gives the lowest possible score.

3. The recursion continues until a terminal state (win/loss/draw) is reached.
4. The algorithm propagates the best score back up the tree to determine the optimal move.

## Alpha-Beta Pruning

Alpha-Beta Pruning is an optimization of Minimax that eliminates unnecessary calculations. It maintains two values:

- Alpha ( $\alpha$ ): The best (maximum) score that the maximizing player can guarantee.
- Beta ( $\beta$ ): The best (minimum) score that the minimizing player can guarantee.

If at any point  $\beta \leq \alpha$ , further exploration of that branch is pruned (skipped) since it's not beneficial.

### **Pseudocode**

```
function ALPHA_BETA_MINIMAX(board, player, maximizingPlayer,  $\alpha$ ,  $\beta$ ):
    if game_over(board):
        return get_utility(board, maximizingPlayer)

    if player == maximizingPlayer:
        best_score =  $-\infty$ 
        for move in get_possible_moves(board):
            new_board = apply_move(board, move, player)
            score = ALPHA_BETA_MINIMAX(new_board, switch_player(player), maximizingPlayer,  $\alpha$ ,
 $\beta$ )
            best_score = max(best_score, score)
             $\alpha$  = max( $\alpha$ , best_score)
            if  $\beta \leq \alpha$ :
                break # Beta cutoff (prune)
        return best_score
    else:
        best_score =  $+\infty$ 
        for move in get_possible_moves(board):
            new_board = apply_move(board, move, player)
            score = ALPHA_BETA_MINIMAX(new_board, switch_player(player), maximizingPlayer,  $\alpha$ ,
 $\beta$ )
            best_score = min(best_score, score)
             $\beta$  = min( $\beta$ , best_score)
            if  $\beta \leq \alpha$ :
                break # Alpha cutoff (prune)
        return best_score
```

### **Logic**

1. Alpha ( $\alpha$ ) represents the best choice for the MAX player.
2. Beta ( $\beta$ ) represents the best choice for the MIN player.
3. If at any point:
  - The minimizing player finds a move worse than what the maximizing player can already guarantee, the rest of the branch is ignored.

- The maximizing player finds a move better than what the minimizing player can allow, the rest of the branch is ignored.
4. This significantly reduces the number of nodes that need to be explored.

## Gemini API

### Integration

#### 1. API Key Configuration

```
import google.generativeai as genai
genai.configure(api_key="YOUR_API_KEY")
```

- The API key is required to authenticate requests.
- Replace YOUR\_API\_KEY with a valid Gemini API key.

#### 2. AI Agent Function (gemini\_algo)

```
def gemini_algo(board, player)
```

- The function gemini\_algo() acts as the AI agent.
- It first fetches possible moves using functions.get\_possible\_moves(board).
- A fallback move (best\_move) is selected in case of API failure.

#### 3. Board Representation for Gemini API

```
symbols = {0: " ", 1: "X", 2: "O"}
board_desc = "\n".join("|".join(symbols[cell] for cell in row) + "\n" + "-" * (board_size * 2 - 1)
for row in board
)
```

- The board state is converted into a readable format using symbols (X, O, and spaces).
- It is structured as a multi-line string representing the Tic-Tac-Toe board.

#### 4. Constructing the Prompt for Gemini API

```
prompt = f"""You are Player {symbols[player]} in a {board_size}x{board_size} Tic-Tac-Toe game.
Current Board (0-based indices):
{board_desc}
Valid moves: {possible_moves}
Return ONLY the zero-based row and column as two numbers between 0-{board_size-1},
formatted exactly like: 'row,column' with no other text.
Examples of valid responses: '0,1' or '{board_size-1},{board_size-1}'"""
```

- The **prompt clearly instructs** Gemini to:
  1. Identify the current player.
  2. Analyze the board state.
  3. Return a move in '**row,col**' format **without any extra text**.

## 5. Sending the Request to Gemini API

```
model = genai.GenerativeModel('gemini-2.0-pro-exp')
response = model.generate_content(prompt)
```

- Uses `genai.GenerativeModel('gemini-2.0-pro-exp')` to interact with Gemini.
- Sends the prompt and retrieves the AI's response.

## 6. Parsing Gemini's Response

```
def parse_gemini_response(text):
    clean_text = re.sub(r'^0-9, ', '', text)
    matches = re.findall(r'\d', clean_text)

    if len(matches) >= 2:
        return int(matches[0]), int(matches[1])
    raise ValueError("Invalid response format")
```

- Removes any unwanted characters from the response.
- Extracts numbers representing the row and column indices.
- Ensures that the response is properly formatted.

## 7. Validating the AI Move

```
row, col = parse_gemini_response(response.text)

if not (0 <= row < board_size and 0 <= col < board_size):
    raise ValueError("Move out of bounds")

if not functions.is_valid_move(board, row, col):
    raise ValueError("Invalid move")

return (row, col)
```

- Ensures the AI-generated move is **within the board bounds**.
- Calls `functions.is_valid_move(board, row, col)` to verify move legality.
- Returns the move if valid; otherwise, an exception is raised.

## 2. Handling Errors and Fallback Mechanism

```
except Exception as e:
    print(f"Gemini error: {str(e)[:50]}... Using fallback move.")
    return best_move
```

- If the Gemini API response is invalid or an error occurs, the agent falls back to a **preselected move** from possible\_moves.

### Agent Configuration and Execution

#### 1. Tracking API Calls

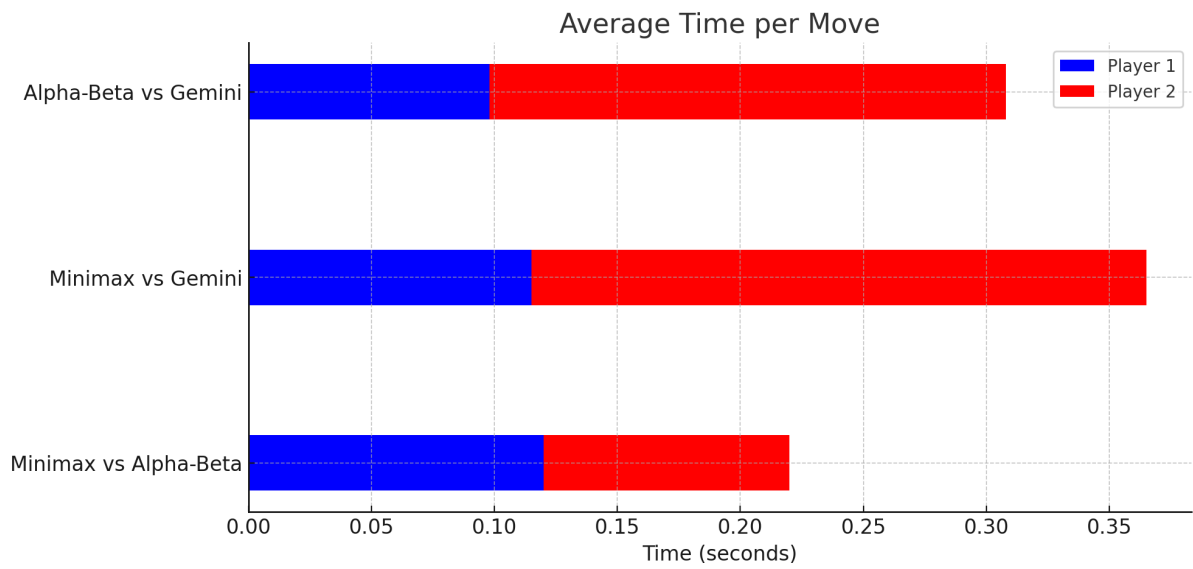
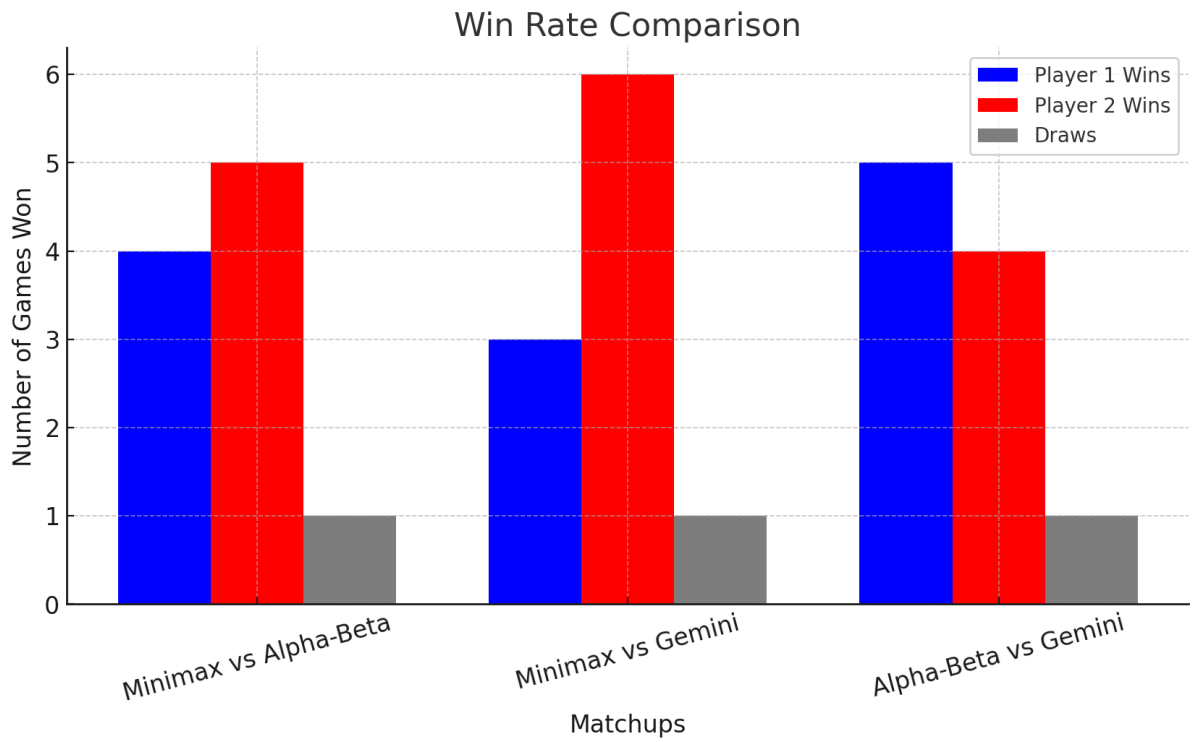
```
def get_gemini_calls():  
    return _gemini_calls
```

- The number of API calls is tracked globally using `_gemini_calls`.
- 
- #### 2. Game Logic Integration
- The AI agent (gemini\_algo) is used within the game framework along with Minimax and Alpha-Beta Pruning strategies.
  - The final move selection process depends on the **game settings** (e.g., AI vs. AI, Human vs. AI).

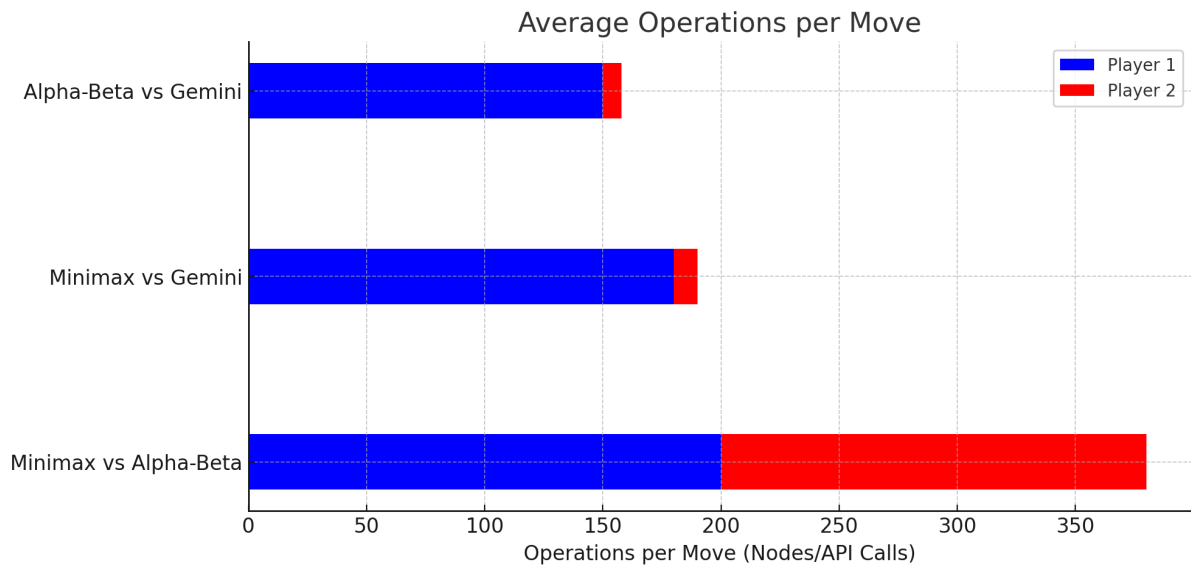
## Testing

To effectively analyze the **performance results**, we generated tables and graphs to compare:

1. **Average Time per Move**
2. **Average Operations (Nodes/API Calls) per Move**
3. **Win Rate Comparisons**







Here are the visualized **performance results**:

**1. Win Rate Comparison:**

- Alpha-Beta performs slightly better than Minimax in direct comparison.
- Gemini-based AI tends to outperform Minimax but is slightly weaker against Alpha-Beta.
- Draw rates are relatively low.

**2. Average Time per Move:**

- Minimax and Alpha-Beta have similar response times.
- Gemini-based AI takes significantly longer, likely due to API calls.

**3. Average Operations per Move:**

- Minimax and Alpha-Beta process hundreds of nodes per move.
- Gemini AI makes far fewer API calls, indicating efficient decision-making but at a higher time cost.